

# Recursive compile time adjoint in C++

23<sup>rd</sup> EuroAD Workshop

Dominic Jones

Siemens PLM, London

`dominic.jones@cd-adapco.com`

11–13 Aug 2020

## Introduction

- The compile time differentiation of a C++ function which only calls built-in functions has been demonstrated in previous talks.
- This talk describes the approach for differentiating a function which may call any nested function.
- Siemen's Simcenter Star-CCM+ simulation software has an implementation of this approach, and is used to differentiate the Spalart Allmaras turbulence model, among other things.

## What's new

### Flat structure (old)

```
hypotenuse(A const &a,  
           B const &b,  
           R &r)  
{  
    auto a2 = a * a;  
    auto b2 = b * b;  
    auto d = a2 + b2;  
    r = sqrt(d);  
}
```

### Nested structure (new)

```
hypotenuse(A const &a,  
           B const &b,  
           R &r)  
{  
    auto d = discriminant(a, b);  
    r = sqrt(d);  
}
```

```
discriminant(A const &a,  
              B const &b)  
{  
    auto a2 = a * a;  
    auto b2 = b * b;  
    return a2 + b2;  
}
```

## Built-in functions

# Hypotenuse

$$r = \sqrt{a^2 + b^2}$$

```
float a = 3;
float b = 4;
float r;

{
    auto d = a*a + b*b;
    r = sqrt(d);
}

std::cout << r << std::endl; // r = 5
```

# Primal of Hypotenuse

$$r = \sqrt{a^2 + b^2}$$

```
float a = 3;
float b = 4;
float r;

auto constexpr mode = DrvMode::PRIMAL;

Drv<mode, float> a_{a}; // input
Drv<mode, float> b_{b}; // input
Drv<mode, float&> r_{r}; // output

{
    auto d = a_*a_ + b_*b_;
    r_ = sqrt(d);
}

std::cout << r << std::endl; // r = 5
```

# Tangents of Hypotenuse

$$\frac{dr}{da}$$

```
float a = 3, a_drv = 1; // w.r.t. 'a'
float b = 4, b_drv = 0;
float      r_drv;

auto constexpr mode = DrvMode::TANGENT;

Drv<mode, float>  a_{a, a_drv};
Drv<mode, float>  b_{b, b_drv};
Drv<mode, float&> r_{r_drv};

{
    auto d = a_*a_ + b_*b_;
    r_ = sqrt(d);
}

std::cout << r_drv << std::endl; // dr/da = 0.6
```

# Tangents of Hypotenuse

$$\frac{dr}{db}$$

```
float a = 3, a_drv = 0;
float b = 4, b_drv = 1; // w.r.t. 'b'
float      r_drv;

auto constexpr mode = DrvMode::TANGENT;

Drv<mode, float>  a_{a, a_drv};
Drv<mode, float>  b_{b, b_drv};
Drv<mode, float&> r_{r_drv};

{
    auto d = a_*a_ + b_*b_;
    r_ = sqrt(d);
}

std::cout << r_drv << std::endl; // dr/db = 0.8
```



# Adjoint of Hypotenuse

$$\begin{bmatrix} \frac{dr}{da} & \frac{dr}{db} \end{bmatrix}^T$$

```
float a = 3, a_drv = 0;
float b = 4, b_drv = 0;
float      r_drv = 1;  // w.r.t. 'r'

auto constexpr mode = DrvMode::ADJOINT;

Drv<mode, float>  a_{a, a_drv};
Drv<mode, float>  b_{b, b_drv};
Drv<mode, float&> r_{r_drv};

{
    auto d = a_*a_ + b_*b_;
    r_ = sqrt(d);
}

std::cout << a_drv << std::endl;  // dr/da = 0.6
std::cout << b_drv << std::endl;  // dr/db = 0.8
```

## User defined functions

## Defining a function

```
struct Discriminant
{
    template<DrvMode mode> // 'float' could be templated, too
    static void
    evaluate(Drv<mode, float> const &a,
            Drv<mode, float> const &b,
            Drv<mode, float&> r)
    {
        auto a2 = a * a;
        auto b2 = b * b;
        r = a2 + b2;
    }
};
```

- All function arguments are treated as differentiable terms, even if in the function body some may be treated passively, e.g.

```
auto a2 = primal(a * a); // i.e. float a2 = a * a;
```

## Calling it via a free function

```
template<typename E0, typename E1>
auto
discriminant(E0 &&e0, E1 &&e1)
->
DrvVariadicNode<GetDrvMode<E0, E1>::value,    // mode
                 decltype(primal(e0 + e1)),    // result
                 Bind_evaluate<Discriminant>,  // operator
                 E0, E1>                      // children
{
    return {std::forward<E0>(e0), std::forward<E1>(e1)};
}
```

## Unified call syntax

```
Drv<mode, float> a_{a, a_drv};  
Drv<mode, float> b_{b, b_drv};  
Drv<mode, float&> r_{r_drv};  
  
{  
    auto d = discriminant(a_, b_);  
    r_ = sqrt(d);  
}
```

- **discriminant** can be used just like any built-in function, such as `sqrt`.
- `a_` or `b_` could be replaced with `a` or `b`, respectively, to treat either passively.

## The variadic expression node

## Unary, binary, ... and *variadic*

- Allied with variadic templates, perfect forwarding and `std::tuple`, unary and binary expression nodes can be generalised into an  $N$ -ary form.
- In principle, all built-in functions, such as `+` `-` `*` `/` and `sin` `cos` `tan` `pow`, could return an  $N$ -ary node.
- A mechanism to distinguish between built-in and user defined functions is still required. This is done with an operator binding, such as `Bind_evaluate<OP>`.

## Perfect forwarding

```
template<class OP, class E0>
struct Unary { Unary(E0 &&e0) {} };
```

```
class Sqrt;
```

```
template<class E0>
auto sqrt(E0 &&e0)
{
    return Unary<Sqrt, E0>(std::forward<E0>(e0));
}
```

```
float const a0{1};
auto a1 = sqrt(a0);           // Unary<Sqrt, float const &>

float b0{1};
auto b1 = sqrt(b0);           // Unary<Sqrt, float &>

auto c1 = sqrt(float{1});     // Unary<Sqrt, float>
```



## Base of all nodes

```
template<DrvMode m, typename R, typename E>
struct DrvExpression
{
    static constexpr auto mode = m;
    using Result = R;
    using Expression = E;
};
```

## Variadic node\*

```
template<DrvMode m,           // mode
        typename R,         // result
        typename OP,        // operator
        typename... EE>     // children
struct DrvVariadicNode
    : DrvExpression<m, R, DrvVariadicNode>
{
    DrvVariadicNode(EE &&... ee) : _ee(ee...) {}

    template<int I> auto const &node() const {
        return std::get<I>(_ee);
    }

    private: std::tuple<EE...> _ee;
};
```

## Variadic node

- The  $N$ -ary node is minimal, providing nothing more than access to its children.
- Children will be associated to the node by value, for in-place expressions or literals, or by reference. This difference is distinguished by perfect forwarding.
- The values and references are held in the `std::tuple` member.

## Evaluating the primal\*

```
template<DrvMode m,
        typename OP, typename R, typename... EE,
        typename AA, int... I>
auto
evaluatePrimal(
    DrvVariadicNode<m, R, OP, EE...> const &expr,
    AA const &aa,                               // primal values: a, b
    std::index_sequence<I...>)
{
    auto constexpr lm = DrvMode::PRIMAL; // local mode
    auto r_pri = R(0);

    OP::evaluate(
        Drv<lm, decltype(DrvVariadicNode::node<I>())::Result>
        {std::get<I>(aa)}...,                // inputs
        Drv<lm, R>{r_pri});                  // output

    return r_pri;
}
```

## Evaluating the primal

- All work relating to differentiable expressions is done by free functions.
- The expression mode,  $m$ , is distinguished from the local mode,  $lm$ , in order to generate the correct instance of the user defined function.
- The arguments for `OP::evaluate` are constructed in-place, generating a new (and independent) differentiation context inside the user defined function.

## Evaluating the adjoint\*

```
template<DrvMode m,
        typename OP, typename R, typename... EE,
        typename AA, typename RHS, int... I>
void
evaluateAdjoint(
    DrvVariadicNode<m, R, OP, EE...> const &expr,
    AA const &aa,                      // primal values: a, b
    AA &aa_drv,                        // adjoint refs:  a_drv, b_drv
    RHS const &rhs,                    // seed value:   r_drv
    std::index_sequence<I...>)
{
    OP::evaluate(
        Drv<m, decltype(DrvVariadicNode::node<I>())::Result>
        {std::get<I>(aa), std::get<I>(aa_drv)}...,
        Drv<m, R&>{rhs});
}
```

## Evaluating the adjoint

- Here, `OP::evaluate` is compiled in adjoint mode, as deduced from the arguments.
- Despite the fact that `OP::evaluate` will have been evaluated in primal mode already, the values of any intermediate results cannot be captured and used here because there is no monolithic expression tree ‘view’.
- The work presented in EuroAD 2019 did have such a view, and so avoided duplicate primal evaluations. But compilation time proved to be too expensive.

## Summary

- The variadic template expression node, combined with perfect forwarding, facilitates generic user defined functions.
- User defined virtual functions can be supported, and so too functions with multiple results.
- Despite being difficult to implement, the added functionality of variadic nodes has proved to be a major advancement in deploying automatic differentiation in Star-CCM+.
- The tool is fully `constexpr` qualified, and with aggressive compiler optimisations performs very well: 2.82x on the harmonic test (331 nodes, 5 inputs); cf. EuroAD 2019.



# Recursive compile time adjoint in C++

23<sup>rd</sup> EuroAD Workshop

Dominic Jones

Siemens PLM, London

`dominic.jones@cd-adapco.com`

11–13 Aug 2020