

Expression tree transforms

For compile-time differentiation

Dominic Jones

`dominic.jones@gmx.co.uk`

March 2018

Differentiate a function *losslessly*

Parity preserving transform

write something like this ...

```
fn(A const &a, B const &b,  
    R &r)  
{  
    auto c0 = 7;  
    auto c1 = 9;  
  
    auto t0 = a * b;    // 1  
    auto t1 = c0 + t0;  // 2  
    auto t2 = c1 + t0;  // 3  
  
    r = t1 / t2;        // 4  
}
```

to implement something like this

```
fn(A const &a, B const &b,  
    R &r)  
{  
    auto c0 = 7;  
    auto c1 = 9;  
  
    auto t0 = a * b;    // 1  
    auto t1 = c0 + t0;  // 2  
    auto t2 = c1 + t0;  // 3  
  
    // somehow add this...  
    t1.d += (1/t2)      * r.d; // 4'  
    t2.d -= (t1/t2^2)   * r.d; // 4'  
    t0.d += t2.d;       // 3'  
    t0.d += t1.d;       // 2'  
    a.d += b * t0.d;    // 1'  
    b.d += a * t0.d;    // 1'  
}
```

Observations


- ➊ Given a pure-functional algorithm, differentiate it
- ➋ Implement the transpose of the chain of derivatives (the adjoint)
- ➌ The required 'extra' code is in the reversed sequence of the original and the data flow is reversed
- ➍ Eager and lazy evaluation: ctor-dtor pairs?

Two hurdles

Dealing with duplicate nodes

Eager evaluation and capture by reference?

```
fn(A const &a, B const &b,  
    R &r)  
{  
    auto c0 = 7;  
    auto c1 = 9;  
  
    auto t0 = a * b;  
    auto t1 = c0 + t0; // l.b.  
    auto t2 = c1 + t0; // r.b.  
  
    r = t1 / t2;        // root  
}
```



Dealing with nested scoping

The *complete* tree, including **cm**, is needed for the transform

```
fn(A const &a, B const &b,  
    R &r)  
{  
    auto c0 = 7;  
    auto c1 = 9;  
  
    auto t0 = a * b;    // 1  
    auto t1 = c0 + t0;  // 2  
    auto t2 = c1 + t0;  // 3  
  
    r = t1 / t2;        // 4  
}
```

```
mul(A const &a, B const &b)  
{  
    auto cm = 1; // lost on return!  
    return cm * a * b;  
}  
  
fn(A const &a, B const &b, R &r)  
{  
    auto c0 = 7;  
    auto c1 = 9;  
  
    auto t0 = mul(a, b);  
    auto t1 = c0 + t0;  
    auto t2 = c1 + t0;  
  
    r = t1 / t2;  
}
```

Dealing with nested scoping

The *complete* tree, including **cm**, is needed for the transform

```
fn(A const &a, B const &b,  
    R &r)  
{  
    auto c0 = 7;  
    auto c1 = 9;  
    auto cm = 1;  
    auto t0 = cm * a * b;  
    auto t1 = c0 + t0;  
    auto t2 = c1 + t0;  
  
    r = t1 / t2;  
}
```

```
fn(A const &a, B const &b,  
    R &r)  
{  
    auto c0 = 3;  
    auto c1 = 4;  
    auto cm = 1;  
    auto t0 = cm * a * b;  
    auto t1 = c0 + t0;  
    auto t2 = c1 + t0;  
  
    // the transform...  
    t1.d += (1/t2)      * r.d;  
    t2.d -= (t1/t2^2) * r.d;  
    t0.d += t2.d;  
    t0.d += t1.d;  
    a.d  += cm * b * t0.d;  
    b.d  += cm * a * t0.d;  
}
```


State of affairs

- ❶ Eager evaluation avoids duplicate branch evaluation - but lazy evaluation will also be needed
- ❷ 'Capture by reference' to keep the tree small - but cannot work with nested scoping
- ❸ 'Capture by value' is too inefficient - the tree will get very large very quickly
- ❹ A monolithic tree, supporting eager and lazy evaluation, of minimal size, and impartial to scoping is required


Only capture terminals? - But how to know if they are already captured?

A 'unique types' approach

Tagging terminals

With all terminal types uniquely tagged, any duplicate branches can be identified

```
fn(A const &a, B const &b,  
    R &r)  
{  
    auto c0 = UQ(7); // unique  
    auto c1 = UQ(9); // unique  
  
    auto t0 = a * b;  
    auto t1 = c0 + t0;  
    auto t2 = c1 + t0;  
  
    r = t1 / t2;  
}
```



Unique

Ugly, but works

```
template<std::size_t ID,
        typename T>
struct Unique
{
    T value;
};

#define UQ(v)          \
Unique<__COUNTER__,    \
      decltype(v)>{v}
```

Interesting, yet useless

```
template<typename T, typename U>
auto constexpr
cmp(T const &t, U const &u)
{
    // Not possible!
    // auto constexpr ta = &t;
    // auto constexpr ua = &u;
    // return ta == ua;
    return &t == &u;
}

{
    auto c0 = 7;
    auto c1 = 9;


    static_assert(cmp(c0, c0));
    static_assert(!cmp(c0, c1));
}
```

Building the components

Bookkeeping

Types and data to percolate up the tree:

- 1 Hierarchically grouped node types
(`Binary<Mul, A, B>`, etc)
- 2 List of constant values
(`c0` and `c1`)
- 3 List of argument addresses
(`a` and `b`)



Hierarchically grouped node types

In the **root** node:

```
// Pad left and right branches to equal depths
```

```
// Merge left and right branches:
```

```
{ {A, B, A, B},  
  {Binary<Mul, A, B>, Binary<Mul, A, B>},  
  {Binary<Add, C0, ...>, Binary<Add, C1, ...>} }
```

```
// Prune duplicates:
```

```
{ {A, B},  
  {Binary<Mul, A, B>},  
  {Binary<Add, C0, ...>, Binary<Add, C1, ...>} }
```

```
// Augment with self:
```

```
{ {A, B},  
  {Binary<Mul, A, B>},  
  {Binary<Add, C0, ...>, Binary<Add, C1, ...>},  
  {Binary<Div, ...C0..., ...C1...>} }
```

Hierarchically grouped node types

- 1 Prune duplicate branches eagerly
- 2 TMP operations exclusively on types
- 3 Relatively low compile-time overhead
- 4 Moderately involved: `resize`, `unique_merge`, `append`
- 5 Using the Brigand library

List of constant values

In the **root** node:

```
// Identify the longer list:
LLC = get_longer<L::LC, R::LC>   i.e. {C0}
SLC = get_shorter<L::LC, R::LC> i.e. {C1}
```

```
// Convert lists to sets:
LSC = brigand::as_set<LLC>
SSC = brigand::as_set<SLC>
```

```
// Merge sets and convert to tuple:
TC = rename<merge_sets<LSC, SSC>, tuple>
TC const tc;   i.e. tuple of constants
```

```
// Difference between tuple and longer list:
S = drop<TC, size<LLC> >   i.e. {C1}
```

```
// Construct 'tc' member: (biggest source of inefficiency!)
tc{tuple_cat(get_longer(l.tc, r.tc),
              select(S{}, get_shorter(l.tc, r.tc)))}
```

List of constant values

- 1 Heavy work done on smaller data
- 2 TMP operations on types and data
- 3 Significant compile-time overhead due to data operations
- 4 Bridging the type-value divide with `select(S, L)`
This should be optimised

select

```
// a naive approach...

template<
    template<class...> class S, class... SS,
    template<class...> class L, class... LL>
S<SS...> select_impl(S<SS...>, L<LL...> l)
{
    // std::get will not work
    return S<SS...>{alt::get<SS>(l)...};
}

template<class S, class L>
S select(S s, L l)
{
    return select_impl(s, l);
}

using S = std::tuple<bool, float>;
using L = std::tuple<char, bool, int, float>;

select(S{}, L{'0', true, 2, 3.0});
```

Possible optimisation:

Given that types are unique and ordered, if any S is found in L, L can be left-cropped

Better to avoid `std::get<T>` altogether

Doing something with all of it

Assigning to the result

```
template<typename T>
struct Result
{
    template<typename Expr>
    void operator=(Expr &expr)
    {
        // depth-grouped node list
        using LN = Expr::LN;

        // list of constants
        using LC = Expr::LC;
        auto &tc = expr.tc;

        // list of argument ptrs
        using LA = Expr::LA;
        auto &ta = expr.ta;

        // Now what?
    }
};
```

Now construct the index arrays to access data:

- 1 build left and right node lists
- 2 build offset arrays for left and right node lists
- 3 build offset array for constants, in order to construct its **dual**

Converting to offsets

```
// Flatten the node list (6 elements)
LN = {A, B,
      Binary<Mul, A, B>,
      Binary<Add, C0, ...>, Binary<Add, C1, ...>,
      Binary<Div, ...C0..., ...C1...>}
```

```
// Offsets of left child nodes ('6' = null marker)
INL = {6, 6, 0, 6, 6, 3}
```

```
// Offsets of right child nodes
INR = {6, 6, 1, 2, 2, 4}
```

```
// With the list of constants...
LC = {Binary<Add, C0, ...>, Binary<Add, C1, ...>}
```

```
// generate node list offsets...
IC = {3, 4}
```

```
// then construct its 'dual' ('2' = null marker)
DC = {2, 2, 2, 0, 1, 2}
```

'Eager and lazy' evaluation

Inside the result assignment operator (pseudocode):

```
// Initialise array/tuple data (6 elements)
```

```
array<double, 6> v = 0;
```

```
// set input from tuple of argument ptrs
```

```
for (i = 0; i != ta.size(); ++i)
```

```
    v[i] = *ta[i];
```

```
// evaluate operators
```

```
for (N: LN)
```

```
    I = offset<N, LN>;
```

```
    N::evaluate<I>(v, tc, INL{}, INR{}, DC{})
```

```
// extract result from 'root' node
```

```
this.r = v[v.size()-1];
```

```
... then the reverse order to differentiate
```

Results

Case 1: nodes: 82, depth: 37, inputs: 2, constants: 12

| Version | compilation | original | auto diff | manual diff |
|------------|-------------|----------|-----------|-------------|
| alt::tuple | 2.2s | 1x | 1.25x | 1.48x |
| std::tuple | 4.5s | 1x | 1.25x | 1.48x |

Case 2: nodes: 331, depth: 25, inputs: 5, constants: 103

| Version | compilation | original | auto diff | manual diff |
|------------|-------------|----------|-----------|-------------|
| alt::tuple | 59s | 1x | 5.7x | 1.9x |
| std::tuple | 27m | 1x | 4.7x | 1.9x |

Conclusion

- ❶ Manipulation the tree is obtained at immense effort
- ❷ Works in the range of exceptionally well (better than hand coded) to acceptably well (better than alternatives)
- ❸ Compile-time features of the language are too limited to use this approach neatly
- ❹ Inlining gives up too readily: `__attribute__((always_inline))` used ubiquitously
- ❺ Not obvious what is going on with `std::tuple`

Footnotes

Learning the ropes

- ① Peter Dimov's "Simple C++ metaprogramming" articles
- ② Time the compilation for large data sets
- ③ An optimised `reverse` function is critical
- ④ Brigand library is an excellent resource

Reflect compile-time literals

Possible

```
// Boost Hana extension...
#ifdef CONFIG_ENABLE_STRING_UDL

auto constexpr x =
    hana::string_c<'a', 'b', 'c'>;

auto constexpr y = "abc"_s;

CONSTANT_CHECK(x == y);
```

Desirable

```
template<auto v>
struct Literal {
    auto constexpr static value = v;
};

auto constexpr x =
    Literal<4.2>;

auto constexpr y = 4.2_f;

static_assert(
    std::is_same_v<decltype(x),
                    decltype(y)>);
```

Reflect compile-time location

Possible

```
123456789 123456789 12345
          10          20
1
2 // main.cpp
3 auto x = 42;
4 decltype(x)* y = &x;
5
reflect          reflect
type            address
```

Desirable

```
123456789 123456789 12345
          10          20
1
2 // main.cpp
3 auto x = 42;
4 auto constexpr y = &x;
5
reflect
location

y = hash(row, column, filename)
  = hash(3, 8, "main.cpp")
```

Function signature pattern

```
struct _1; struct _2;
template<template<class...> class, class...> struct bind;

// Apply a function taking two args to a list...
template<class L, class P> struct apply;

template<template<class...> class L, class T, class... Ts,
        template<class...> class F>
struct apply<L<T, Ts...>, bind<F, _1, _2> >
{
    using type = L<F<T, Ts>...>;
};

// A function taking two args...
template<class T, class U>
using cmp = bool_t<(sizeof(T) > sizeof(U))>;

// evaluate 'cmp' over a list
using L = list<int, bool, char, double>;
using R = apply<L, bind<cmp, _1, _2> >::type;
```