

Sequential Processing in Nature, 'Anything but' in Scientific Computation

an observation


Dominic Jones

Netherhall House, London

January 2018

Nature - 'It should just work'


Finding the global minimum



Folding occurs in microseconds, but cannot be predicted

Abnormal folding

Deformed proteins cannot be mended




In brief

- 1 From a chain of amino acids a specific highly complex shape is required
- 2 There is no apparent error correction in the process (or very little)
- 3 It is usually successful
- 4 When it isn't, Creutzfeldt-Jakob disease, Parkinson's, Alzheimer's, etc


Scientific computation - ' $T = \text{abs}(T)$ '

Attempts at prediction and projection

Numerical approximation




Local sensitivities



Iterative, discretised and linearised algorithms

All starts with a graph (mesh)

Resolve change at the small scales



Mesh implies graph, implies matrix, implies matrix inversion

Matrix inversion: scaling up

```

void inverse(int n, float a[][2*n_max])
{
    for (int i = 0; i < n; i++)
        for (int j = n; j < 2*n; j++)
            a[i][j] = (i == j-n? 1: 0);


    for (int i = 0; i < n; i++) {                // linear
        float aii = a[i][i];
        for (int j = i; j < 2*n; j++)
            a[i][j] = a[i][j] / aii;

        for (int j = 0; j < n; j++) {            // quadratic!
            if (i != j) {
                float aji = a[j][i];
                for (int k = 0; k < 2*n; k++)    // cubic!!
                    a[j][k] = a[j][k] - aji * a[i][k];
            }
        }
    }
}

```


Computation: slow but compact

Geometry and mesh



$\begin{matrix} y \\ | \\ x \end{matrix}$


Connectivity graph



$\begin{matrix} y \\ | \\ x \end{matrix}$

Topology Representation

Symmetric, sparse (48 non-zeros), irregular



Sparsity and Indirection

- Dense storage: 196 values, 24% efficient
- Direct access to values

```
float[14][14] A;  
A[2][3] = 3.142;
```

- Compressed row storage: 111 values, 56% efficient
- Requires indirection to access values (10x slower)

```
int[15] IA = [0, 3, 7, 10, ...];  
int[48] JA = [0, 1, 8, 0, 1, 2, 6, 1, 2, 3, ...];  
float[48] A;  
  
A[JA[IA[2]+2]] = 3.142; // i.e. A[2][3] = 3.142;
```

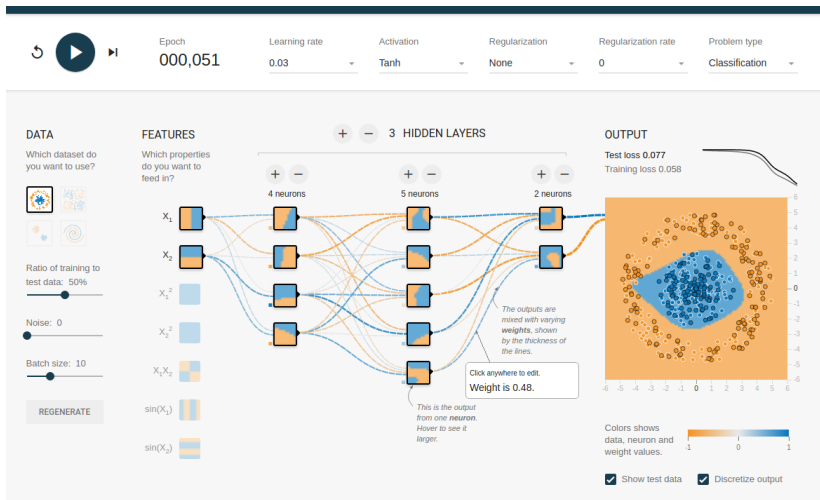
In brief

- 1 Many scientific algorithms have graphs and matrices, and perform matrix inversion at their core
- 2 They are often woefully inefficient due to the requirement to use compressed storage of irregular data
- 3 The compromise on efficiency is to leverage compact memory footprint
- 4 In the process of linearising, much of the 'beauty' of underlying the mathematics gives way to crude approximation

Bridging the Nature — Computation dichotomy

Artificial Neural Networks

TensorFlow demo



Properties of ANN

- 1 Weights are associated with nodes and inter-nodal connections
- 2 Processes are relatively straight forward - matrix-vector products and local reductions
- 3 Its kernel is essentially plastic - number of hidden layers and number of nodes on each layer govern its form
- 4 It has somewhat of the 'it should just work' essence - a result will always be produced

Where the analogy breaks down

- ❶ *These* nodes and these connections are trained for *this* problem
- ❷ Training is complicated - requires derivative of every output with respect to every weight
- ❸ Improvement of the weights requires vast resources relative to evaluating the actual problem

In brief

- 1 Somehow, the need for feedback (or at least how it is presently done) appears to lack a natural analogy
- 2 For protein folding, it is as though all possible solutions (and so the right one) are ‘known’ at once
- 3 This is somewhat characteristic of quantum computing, and moreover, its solution is *probably* correct

Incidentally...

Incidentally...

Transpose is like dependency tracing

- 1 Given a sequence of operations: $X(D)$, $Q(X)$, $L(Q)$, the Jacobian of Q w.r.t. X can be written as

$$J_Q = \begin{bmatrix} \frac{\partial Q_1}{\partial X_1} & \cdots & \frac{\partial Q_1}{\partial X_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial Q_m}{\partial X_1} & \cdots & \frac{\partial Q_m}{\partial X_n} \end{bmatrix} = \frac{dQ}{dX}$$

- 2 The derivative of the whole system is the chain of the Jacobians of each operation, giving

$$\frac{dL}{dD} = \frac{dL}{dQ} \frac{dQ}{dX} \frac{dX}{dD}$$

Transpose is like dependency tracing

- 1 The derivative of the system could be built up by inputting *tangents* of J_X

$$\frac{dL}{dD_i} = \frac{dL}{dQ} \frac{dQ}{dX} \frac{dX}{dD_i}$$

- 2 But by turning the tangents approach “inside-out”, taking the transpose of the derivative of the system for a particular *gradient* of J_L

$$\frac{dL_j}{dD}^T = \frac{dX^T}{dD} \frac{dQ^T}{dX} \frac{dL_j}{dQ}^T$$

offers an efficient way of relating an output to all of its inputs.

Give the output a handle on its history


```

auto eval(A const &a,
           B const &b)
{
    auto c0 = 3;
    auto c1 = 4;

    auto t0 = a * b;
    auto t1 = c0 + t0;
    auto t2 = c1 + t0;
    auto r  = t1 / t2;

    return r;
}

```



Instead of computing values, construct expression trees (graphs!)

'Adjoint' differentiation

Apply the chain rule and transpose...

```
t0 = a * b      // 1
t1 = c0 + t0    // 2
t2 = c1 + t0    // 3
r  = t1 / t2    // 4
```

```
// 4
t1_d += (1 / t2) * r_d
t2_d -= (t1 / t2^2) * r_d
```


```
// 3
t0_d += t2_d
```

```
// 2
t0_d += t1_d
```

```
// 1
a_d += b * t0_d
b_d += a * t0_d
```

Directed design

Move surface **in** or **out** to reduce drag



In brief

- 1 There is some analogy with entanglement and wave function collapsing in classical computation
- 2 It gives all solutions in one evaluation
- 3 But, everything really is back-to-front! (and implementing it is a nightmare!)